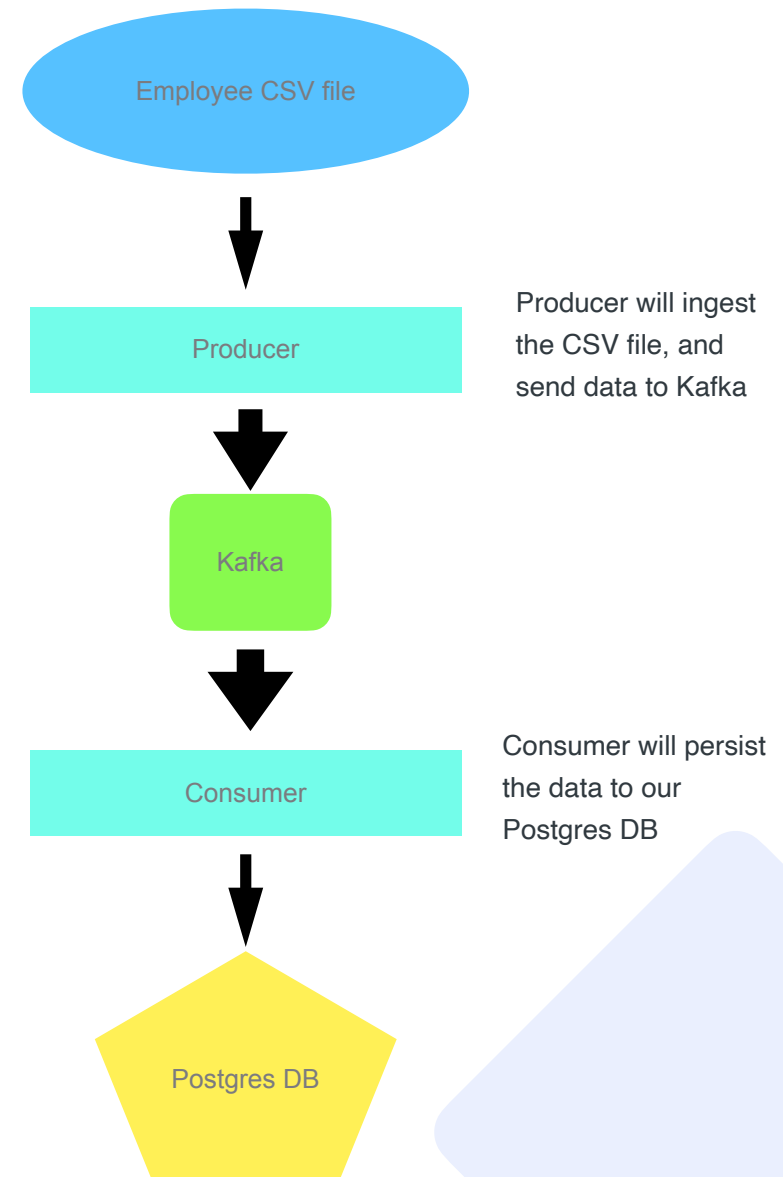


Kafka Project 1



Kafka Project 1: ETL pipeline

- This project will emphasize on creating a data pipeline in which CSV files are being generated on a regular basis, and then data is being cleaned, transformed and then later persisted into some storage for data analytics.
- To calculate the total salary of the departments.



Data Producer

DataProducer Responsibilities - (Extract + Transform)

- Ingest Employee_Salaries.csv file (will be in resources folder)
- Perform these transformations -
 - Ingest only these Departments -
 - ECC
 - CIT
 - EMS
 - Round off the Salary to lower number
 - Employees hired after 2010 (≥ 2010)
- Send this data to Kafka

Data Consumer

DataConsumer Responsibilities - (Load)

- Ingest the data into the **Department_Employee** Table which will have this schema -
 - department_division: varchar
 - position_title: varchar
 - hire_date: Date
 - salary: int32
- With every message, also update the total salary given by each department. Schema for the department table -
 - department: varchar
 - total_salary: int64

Table Schema

```
CREATE TABLE department_employee(  
  department VARCHAR(200),  
  department_division VARCHAR(200),  
  position_title VARCHAR(200),  
  hire_date DATE,  
  salary decimal  
);
```

```
CREATE TABLE  
public.department_employee_salary (  
  department varchar(200) NOT NULL,  
  total_salary int4 NULL,  
  
  CONSTRAINT department_employee_salary_pk PRIMARY KEY (department)  
);
```

helper code


```
# update statement to update the table

cur.execute(
    "insert into department_employee_salary (department,total_salary)
    values (%s,%s) on conflict(department) do update set total_salary =
    department_employee_salary.total_salary + %s      ",
    (e.department, int(float(e.salary)), int(float(e.salary))))
```

Success Criteria

You should get the same results.

It's OK if you don't get exactly the same results.


department_employee_salary  Enter a SQL expression			
	department	total_salary	
1	EMS	3,779,570	
2	CIT	9,102,142	
3	ECC	2,042,698	

Kafka Project 2



Kafka Project 2: Changing Data Capture

- For this project, we will create a data tunnel to **maintain a sync** between 2 tables in two different databases.
- The stream processing would be responsible for:
 - snapshot processing => Reading all data from db1 => send asynchronously to db2
 - stream processing => Once a snapshot has been synchronized, stream processing will keep looking for the new
- Insert/Delete/Update => any changes on Emp_A should be reflected in Emp_B **< 1 sec.**



Two approaches

Polling on the tables -

With this approach we can create a python program which will start scanning the data from the head of the table till the last row, with every row being scanned will be sent to a streaming pipeline. This program would need to keep the track of the last offset/row which has been consumed, otherwise on the program crash this program would need to start all over again. The only problem with this approach is - this will take care of inserts, but not updates/deletes.

Use this approach!!!!



SQL Triggering (Postgres Triggers) -

Triggers + Functions

- Using Postgres triggers we can call an SQL function on every row insert/update/delete.
- The SQL function in this case should insert all the affected rows plus actions (insert, update or delete) into a new table. This new table will act as a CDC table.
- Producer will start scanning the data from the head of the CDC table till the last row, with every row being scanned will be sent to a streaming pipeline (Kafka). This program would need to keep the track of the last offset/row which has been consumed.
 - Otherwise we will perform full scan every time, and on the program crash this program would need to start all over again, which is very inefficient.
- Now this action should be passed as well as with the other information to the Kafka, and the consumers consuming this data will update all the syncs.

Steps

- 1, Modify the Docker compose file to have two databases, both port-forward to different ports.
- 2, Use the same schema for employee table and add this table in both databases
 - `CREATE TABLE employees(emp_id SERIAL, first_name VARCHAR(100), last_name VARCHAR(100), dob DATE, city VARCHAR(100), salary INT);`
- 3, Add PSQL functions and triggers on the employee_A table, which will insert the rows to the new cdc table.
 - `CREATE TABLE emp_cdc(emp_id SERIAL, first_name VARCHAR(100), last_name VARCHAR(100), dob DATE, city VARCHAR(100), salary INT, action VARCHAR(100));`
- 4, Modify the producer and consumer code (depending on your won configuration) to scan the employee_cdc table, and send the records to the topic; and to consume the data and update the employee_B table based on the action.

Some good practices

- There is no best solution to the problem
- Try implementing everything in an OOP manner to practice
- Always try and catch all exception you might have, especially in a production environment.
- Maintain a good coding style (ifUsingCamelToeNotation_then_do_not_switch_to_another_style)
- Some optional things you can try (absolutely no bonus)
 - Use a git repo for version control
 - Scale up/down your instances
 - Create a DLQ
 - Use Offset Explorer
 - Can use AirFlow to schedule producer.

Optional: DLQ

- Set up a DLQ topic in your docker-compose.yml file, or through admin.py (check confluent official documentations)
- You can implement your custom logic, i.e.
 - all employees should be born after 2007
 - salary should be greater than 10000
 - no negative emp id
 - ...
- See if you can intentionally create some bad actions, and if you are able to get everything to work.